

CMPS 181, Spring 2016, Project 2

Due Friday, May 6, 11:55pm on eCommons

Course Project 2 Description

Introduction

In this project, you will continue implementing the record-based file manager (RBFM). Once you have finished implementing that, you will build a relation manager (RM) on top of the basic paged file system. The RM manager should meet the following basic requirements.

Basic Requirements (100 points)

Finish RBFM

You need to finish the implementation of the Record-Based File Manager (RBFM) that you started in Project 1. As the basis for Project 2, you may use your solution to Project 1, or TA Coy Humphrey's Project 1 solution, or a combination of the two. Specifically, you should finish implementing the following methods: **RecordBasedFileManager::deleteRecord()**, **RecordBasedFileManager::updateRecord()**, **RecordBasedFileManager::readAttribute()**, and **RecordBasedFileManager::scan()**. Please refer to the description of Project for the explanation of these methods, and look at the file **rbfm.h** for their signatures. Once you finish the implementation of those methods and test them, you will use them to implement the Relation Manager; please look at the file **rm.h**.

Implement the Relation Manager

The RelationManager class is responsible for managing the database tables. It handles the creation and deletion of tables. It also handles basic relational operations performed on top of a table (e.g., insertion and deletion of tuples).

Catalog

Create a Catalog to hold information about your database. This includes at least the following:

- Table information (e.g., table-name, table-id, etc.).
- For each table, the columns in that table. For each of these columns: the column name, type, length, and position.
- The name of the Record-Based File in which the data corresponding to each table is stored.

It is mandatory to store Catalog information by using the RBF layer functions. You should create the Catalog's tables, and populate them the first time your database is initialized (when the method `createCatalog()` is called). Once the Catalog's tables and columns tables have been created, they should be persisted to disk. Please use the following names and types for the catalog tables and columns. You can add additional attributes if you want to/need to do so. However, please do not change the specified names of these two tables or their attribute names.

```

Tables (table-id:int, table-name:varchar(50), file-name:varchar(50))

Columns(table-id:int, column-name:varchar(50), column-type:int, column-length:int,
column-position:int)

```

An example of the records that should be in these two tables after creating a table named "Employee" is:

```

Tables
(1, "Tables", "Tables")
(2, "Columns", "Columns")
(3, "Employee", "Employee")

Columns
(1, "table-id", TypeInt, 4, 1)
(1, "table-name", TypeVarChar, 50, 2)
(1, "file-name", TypeVarChar, 50, 3)
(2, "table-id", TypeInt, 4, 1)
(2, "column-name", TypeVarChar, 50, 2)
(2, "column-type", TypeInt, 4, 3)
(2, "column-length", TypeInt, 4, 4)
(2, "column-position", TypeInt, 4, 5)
(3, "empname", TypeVarChar, 30, 1)
(3, "age", TypeInt, 4, 2)
(3, "height", TypeReal, 4, 3)
(3, "salary", TypeInt, 4, 4)

```

Note that TypeInt, TypeVarChar, and TypeReal are the enum types used to represent an attribute type, which are defined in the rbfm.h file. For the file-name attribute, you may use your own naming conventions if you want (e.g., "Tables.tbl" for the Tables table). In the above example, the last row of the Tables table states that the table-id of the "Employee" table is 3, its name is "Employee", and the RBF file associated with this table is "Employee". The last line of the Columns table shows the information about the "salary" column of the "Employee" table. For this column, the table-id is 3, the column-name is "salary", the type of this column is Int, its length is 4, and the position of this column is 4.

The Catalog tables (Tables and Columns) should be created when the method createCatalog() is called. All other subsequent invocations of your database should use the already created Catalog tables. The Catalog tables will be deleted when deleteCatalog() is called. Between the createCatalog() and deleteCatalog() methods, the Catalog should be persistent on disk. Please note that users should be able to query your Catalog tables just as they can query any other tables, but they should not be allowed to modify Catalog table content through the RM API. Modifications of Catalog tables should only be possible using internal calls (for instance, when creating a user table or deleting a user table). To implement this, you may need to have a "system" vs. "user" flag in your Tables Catalog, so that you can distinguish "system tables" from "user tables", and make it illegal to do anything but read "system tables" through the RM layer. If you follow this design, then you may need to add one more column to two schemas defined above to represent this flag.

RelationManager class

The following provides more information about the RelationManager class. Your program should create exactly one instance of this class, and all requests for the RM component should be directed to that instance. The public methods of the class declaration are shown first, followed by descriptions of the methods. The last two methods in the class declaration are the constructor and destructor methods; they are not explained further. **Note:** your tuple-oriented file system must create a relation manager (RM) that initializes the Catalog information that you need to store. It also internally creates a Record-Based File Manager using the implementation from Project 1.

```

class RelationManager
{
public:
    RC createCatalog();

```

```

RC deleteCatalog();

RC createTable(const string &tableName, const vector<Attribute> &attrs);

RC deleteTable(const string &tableName);

RC getAttributes(const string &tableName, vector<Attribute> &attrs);

RC insertTuple(const string &tableName, const void *data, RID &rid);

RC deleteTuple(const string &tableName, const RID &rid);

RC updateTuple(const string &tableName, const void *data, const RID &rid);

RC readTuple(const string &tableName, const RID &rid, void *data);

// mainly for debugging
// Print a tuple that is passed to this utility method.
RC printTuple(const vector<Attribute> &attrs, const void *data);

RC readAttribute(const string &tableName, const RID &rid, const string &attributeName,
void *data);

// scan returns an iterator to allow the caller to go through the results one by one.
RC scan(const string &tableName,
const string &conditionAttribute,
const CompOp compOp,           // comparison type such as "<" and "="
const void *value,             // used in the comparison
const vector<string> &attributeNames, // a list of projected attributes
RM_ScanIterator &rm_ScanIterator);

protected:
    RelationManager();
    ~RelationManager();
};

```

RC createCatalog()

This method creates two system catalog tables - Tables and Columns. If they already exist, return an error. The actual files for these two tables should be created, and tuples describing these two system catalog tables themselves should be inserted into these tables, as shown earlier in the above Catalog section of this document.

RC deleteCatalog()

This method deletes the system catalog tables. The actual files for these two tables should be deleted. Return an error if the system catalog does not exist.

RC createTable(const string &tableName, const vector<Attribute> &attrs)

This method creates a table called tableName with a vector of attributes (attrs). The actual RBF file for this table should be created. Return an error if a table called tableName already exists.

RC deleteTable(const string &tableName);

This method deletes the table that has the specified tableName. The actual RBF file for this table should be deleted. Return an error if a table called tableName does not exist.

RC getAttributes(const string &tableName, vector<Attribute> &attrs);

This method gets the attributes (attrs) of the table called tableName by looking in the catalog tables. Return an error if a table called tableName does not exist.

RC insertTuple(const string &tableName, const void *data, RID &rid);

This method inserts a tuple into the table called tableName. Return an error if a table called tableName does not exist. You can assume that the input is always correct and free of error. That is, you do not need to check if the input tuple has the right number of attributes and/or if the attribute types match. Since there can be NULL values in one or more attributes, the first part in *data contains n bytes to pass the null information about each attributes. For details, see insertRecord() in the description of Project 1.

RC deleteTuple(const string &tableName, const RID &rid);

This method deletes a tuple with the specified rid. Each time a tuple is deleted, you will need to compact the underlying page. That is, keep the free space together in the middle of the page -- the slot table should be at one end of the page, the record data area should be at the other end, and the free space should be in the middle.

Return an error if a table called tableName does not exist. Also return a (different) error if there is no tuple with the specified rid.

RC updateTuple(const string &tableName, const void *data, const RID &rid);

This method updates a tuple identified by the specified rid. **Note:** if the tuple grows (i.e., the size of the tuple increases) and there is no space in the page to store the tuple (after the update), then the tuple is migrated to a new page with enough free space. Since you will implement an index structure (e.g., B-tree) in Project 3, tuples are identified by their rids and when they migrate, you must leave a forwarding address behind identifying the new location of the tuple. Also, each time a tuple is updated and becomes smaller, you need to compact the underlying page. That is, keep the free space in the middle of the page -- the slot table should be at one end of the page, the tuple data area should be at the other end, and the free space should be in the middle. Again, the structure for *data is the same as for insertRecord().

Return an error if a table called tableName does not exist. Also return a (different) error if there is no tuple with the specified rid.

RC readTuple(const string &tableName, const RID &rid, void *data);

This method reads a tuple identified by a given rid. The structure for *data is the same as for insertRecord(). Return an error if a table called tableName does not exist. Also return a (different) error if there is no tuple with the specified rid.

RC printTuple(const vector<Attribute> &attrs, const void *data);

This method mainly exists for debugging purposes. This method prints the tuple whose data is passed into this method. The structure for *data is the same as for insertRecord(). For details, refer to printRecord() in the Project 1 description.

RC readAttribute(const string &tableName, const RID &rid, const string &attributeName, void *data);

This method mainly exists for debugging purposes. This method reads a specific attribute of a tuple identified by a given rid. The structure for *data is the same as for insertRecord(). That is, a null-indicator will be placed in the beginning of *data. However, for this function, since it returns a value for just one attribute, exactly one byte of null-indicators should be returned, not a set of the null-indicators for all of the tuple's attributes.

Return an error if a table called tableName does not exist. Also return a (different) error if there is no tuple with the specified rid.

RC scan(const string &tableName, const string &conditionAttribute, const CompOp compOp, const void *value, const vector<string> &attributeNames, RM_ScanIterator &rm_ScanIterator);

This method scans the table called tableName. That is, it sequentially reads all of the tuples in the table. This method returns an iterator called rm_ScanIterator to allow the caller to go through the records in the table one by one. A scan has a filter condition associated with it, e.g., it consists of a list of attributes to project, as well as a predicate on an attribute (such as "Sal > 40000"). **Note:** the RBFM_ScanIterator should not cache the scan result in memory. Your code should only be looking at one (or a few) page(s) of data at a time when getNextTuple() is called. In this project, let the OS do all the memory management work for you.

RM_ScanIterator Class

The RM_ScanIterator class is a class that represents an iterator that is used to go through the tuples in the table one by one. The way to use this iterator is as follows:

```
RM_ScanIterator rmsi;

// At this moment, do not execute scan and cache the results in the memory. Just
initialize the scan operator
rc = rm->scan(tableName, conditionAttribute, compOp, value, attributes, rmsi);

while(rmsi.getNextTuple(rid, returnedData) != RM_EOF){
    // fetch one tuple at a time and process the data;
}
rmsi.close();
```

The public methods of this class are shown next. The first two methods in the class declaration are the constructor and destructor methods, and these methods are not explained further.

```
class RM_ScanIterator {
public:
    RM_ScanIterator();
    ~RM_ScanIterator();

    // "data" follows the same format as RelationManager::insertTuple()
    RC getNextTuple(RID &rid, void *data);
    RC close();
};
```

RC getNextTuple(RID &rid, void *data);

This method is used to get the next tuple from the scanned table. It returns RM_EOF when all tuples have been scanned. Note that the structure for *data is the same as for insertRecord(). That is, a null-indicator will be placed in the beginning of *data. However, for this function, since it returns a value for the attributes that are specified in the scan method, the corresponding byte(s) for these attributes of null-indicators should be returned, not the null-indicators for the entire set of table attributes.

RC close();

This method is used to close the iterator.

Design Assumptions

You can make the following simplifying assumptions when implementing PFM, RBFM, and RM (including ScanIterator):

1. The size of one tuple cannot exceed the size of a page. That is, an empty page can always hold at least one tuple. If an incoming tuple is too big for an empty page, or a tuple grows too much to fit alone on a page, return an error. (However, if a page has two or more tuples, then one of the tuples could be updated and grow in size, no longer fitting on that shared page.)
2. A table maps to a single file, and a file contains only one table.

Explanation

The commands listed above are by no means complete, but they do capture the essence of a tuple-oriented file system.

You have a lot of freedom in designing your specific algorithms and building your system. You should spend a significant amount of time in coming up with a design for your system before you start coding.

Grading will be based on the correctness of the implementation, unless your code takes an exceedingly long time to execute a test case. Most test cases can be executed within a few seconds.